

```
In [1]: # Import required packages
import numpy as np
import pandas as pd
import cv2
import matplotlib.pyplot as plt
from sklearn.metrics import classification_report
from sklearn.linear_model import LogisticRegression
import tensorflow as tf
import time

import os
from keras.models import Model
from keras.optimizers import Adam
from keras.applications import vgg16, vgg19, resnet, inception_v3, MobileNetV3Small, n
from keras.preprocessing.image import ImageDataGenerator
from keras.callbacks import ModelCheckpoint, EarlyStopping
from keras.layers import Dense, Dropout, Flatten
from keras.utils import to_categorical
from livelossplot.inputs.keras import PlotLossesCallback
from sklearn.metrics import accuracy_score
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report
from sklearn.svm import SVC
from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import RandomForestRegressor
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import GradientBoostingClassifier
import seaborn as sns
from sklearn.metrics import confusion_matrix
from datetime import datetime as dt

from sklearn.model_selection import KFold

import warnings
warnings.filterwarnings("ignore")
```

1. Load the datasets

For the project, we provide a training set with 50000 images in the directory

`../data/images/` with:

- noisy labels for all images provided in `../data/noisy_label.csv` ;
- clean labels for the first 10000 images provided in `../data/clean_labels.csv` .

```
In [2]: # [DO NOT MODIFY THIS CELL]

# Load the images
n_img = 50000
n_noisy = 40000
n_clean_noisy = n_img - n_noisy
imgs = np.empty((n_img,32,32,3))
for i in range(n_img):
    img_fn = f'../data/images/{i+1:05d}.png'
    imgs[i,:,:,:]=cv2.cvtColor(cv2.imread(img_fn),cv2.COLOR_BGR2RGB)
```

```
# Load the Labels
clean_labels = np.genfromtxt('../data/clean_labels.csv', delimiter=',', dtype="int8")
noisy_labels = np.genfromtxt('../data/noisy_labels.csv', delimiter=',', dtype="int8")
```

For illustration, we present a small subset (of size 8) of the images with their clean and noisy labels in `clean_noisy_trainset`. You are encouraged to explore more characteristics of the label noises on the whole dataset.

In [3]: # [DO NOT MODIFY THIS CELL]

```
fig = plt.figure()

ax1 = fig.add_subplot(2,4,1)
ax1.imshow(imgs[0]/255)
ax2 = fig.add_subplot(2,4,2)
ax2.imshow(imgs[1]/255)
ax3 = fig.add_subplot(2,4,3)
ax3.imshow(imgs[2]/255)
ax4 = fig.add_subplot(2,4,4)
ax4.imshow(imgs[3]/255)
ax1 = fig.add_subplot(2,4,5)
ax1.imshow(imgs[4]/255)
ax2 = fig.add_subplot(2,4,6)
ax2.imshow(imgs[5]/255)
ax3 = fig.add_subplot(2,4,7)
ax3.imshow(imgs[6]/255)
ax4 = fig.add_subplot(2,4,8)
ax4.imshow(imgs[7]/255)

# The class-label correspondence
classes = ('plane', 'car', 'bird', 'cat',
           'deer', 'dog', 'frog', 'horse', 'ship', 'truck')

# print clean labels
print('Clean labels:')
print(' '.join('%5s' % classes[clean_labels[j]] for j in range(8)))
# print noisy labels
print('Noisy labels:')
print(' '.join('%5s' % classes[noisy_labels[j]] for j in range(8)))
```

```
Clean labels:
    frog truck truck deer    car    car   bird horse
Noisy labels:
    cat    dog truck  frog    dog    ship   bird  deer
```



2. The predictive model

We consider a baseline model directly on the noisy dataset without any label corrections. RGB histogram features are extracted to fit a logistic regression model.

2.1. Baseline Model

```
In [4]: # [DO NOT MODIFY THIS CELL]
# RGB histogram dataset construction
no_bins = 6
bins = np.linspace(0,255,no_bins) # the range of the rgb histogram
target_vec = np.empty(n_img)
feature_mtx = np.empty((n_img,3*(len(bins)-1)))
i = 0
for i in range(n_img):
    # The target vector consists of noisy labels
    target_vec[i] = noisy_labels[i]

    # Use the numbers of pixels in each bin for all three channels as the features
    feature1 = np.histogram(imgs[i][:,:,0],bins=bins)[0]
    feature2 = np.histogram(imgs[i][:,:,1],bins=bins)[0]
    feature3 = np.histogram(imgs[i][:,:,2],bins=bins)[0]

    # Concatenate three features
    feature_mtx[i,:] = np.concatenate((feature1, feature2, feature3), axis=None)
    i += 1
```

```
In [5]: # [DO NOT MODIFY THIS CELL]
# Train a Logistic regression model
clf = LogisticRegression(random_state=0).fit(feature_mtx, target_vec)
```

For the convenience of evaluation, we write the following function `predictive_model` that

does the label prediction. **For your predictive model, feel free to modify the function, but make sure the function takes an RGB image of numpy.array format with dimension $32 \times 32 \times 3$ as input, and returns one single label as output.**

```
In [6]: # [DO NOT MODIFY THIS CELL]
def baseline_model(image):
    ...
    This is the baseline predictive model that takes in the image and returns a label
    ...
    feature1 = np.histogram(image[:, :, 0], bins=bins)[0]
    feature2 = np.histogram(image[:, :, 1], bins=bins)[0]
    feature3 = np.histogram(image[:, :, 2], bins=bins)[0]
    feature = np.concatenate((feature1, feature2, feature3), axis=None).reshape(1, -1)
    return clf.predict(feature)
```

```
In [7]: fig = plt.figure()

ax1 = fig.add_subplot(2,4,1)
ax1.imshow(imgs[0]/255)
ax2 = fig.add_subplot(2,4,2)
ax2.imshow(imgs[1]/255)
ax3 = fig.add_subplot(2,4,3)
ax3.imshow(imgs[2]/255)
ax4 = fig.add_subplot(2,4,4)
ax4.imshow(imgs[3]/255)
ax1 = fig.add_subplot(2,4,5)
ax1.imshow(imgs[4]/255)
ax2 = fig.add_subplot(2,4,6)
ax2.imshow(imgs[5]/255)
ax3 = fig.add_subplot(2,4,7)
ax3.imshow(imgs[6]/255)
ax4 = fig.add_subplot(2,4,8)
ax4.imshow(imgs[7]/255)

print('Clean labels:')
print(' '.join('%5s' % classes[clean_labels[j]] for j in range(8)))
print('Predicted baseline labels:')
print(' '.join('%5s' % classes[int(baseline_model(imgs[j])[0])] for j in range(8)))
```

Clean labels:
frog truck truck deer car car bird horse
Predicted baseline labels:
frog ship truck frog ship cat deer horse



2.2. Model I

- Loading images for pre-processing using the pre-trained inception V3 model. The dataset is split into train and test dataset using `train_test_split`

```
In [9]: imgs_load = np.empty((n_img,75,75,3))
for i in range(n_img):
    img_fn = f'../data/images/{i+1:05d}.png'
    imgs_load[i,:,:,:]=cv2.resize(cv2.imread(img_fn),(75,75),interpolation=cv2.INTER_CUBIC)
    imgs_incep = inception_v3.preprocess_input(imgs_load)
```

```
In [7]: # Split the data into training set (75%) and test set (25%)
img_incep_train, img_incep_test, noisy_labels_train, noisy_labels_test = train_test_
```

- We load the pre-trained model (Inception_v3) weights and include layers of fully connected network for training the model. Adam optimizer with `learning_rate=0.001` is used.

```
In [15]: # Load inception_v3 model
def create_model_incep(input_shape, n_classes, optimizer, fine_tune):
    """
    Compiles a model integrated with inception_v3 pretrained layers

    input_shape: tuple - the shape of input images (width, height, channels)
    n_classes: int - number of classes for the output layer
    optimizer: string - instantiated optimizer to use for training. Defaults to 'RMSProp'
    fine_tune: int - The number of pre-trained layers to unfreeze.
                  If set to 0, all pretrained layers will freeze during training
    """
    # Pretrained convolutional Layers are loaded using the Imagenet weights.
    # Include_top is set to False, in order to exclude the model's fully-connected Lay
```

```

conv_base = inception_v3.InceptionV3(include_top=False,
                                      weights='imagenet', input_shape=input_shape)

# Defines how many layers to freeze during training.
# Layers in the convolutional base are switched from trainable to non-trainable
# depending on the size of the fine-tuning parameter.
# If the arg fine_tune is set to 0, all pre-trained layers will be frozen and left
# otherwise, the last n layers will be made available for training.
if fine_tune > 0:
    for layer in conv_base.layers[:-fine_tune]:
        layer.trainable = False
else:
    for layer in conv_base.layers:
        layer.trainable = False

# Create a new 'top' of the model (i.e. fully-connected layers)
# by grabbing the conv_base outputs and flattening them.
# This is 'bootstrapping' a new top_model onto the pretrained layers.
top_model = conv_base.output
top_model = Flatten()(top_model)
top_model = Dense(n_classes*8, activation='relu')(top_model)
top_model = Dense(n_classes*4, activation='relu')(top_model)
top_model = Dropout(0.2)(top_model)
output_layer = Dense(n_classes, activation='softmax')(top_model)

# Group the convolutional base and new fully-connected layers into a Model object.
model = Model(inputs=conv_base.input, outputs=output_layer)

# Compiles the model for training.
model.compile(optimizer=optimizer,
              loss='categorical_crossentropy',
              metrics=['accuracy'])

return model

input_shape = (75, 75, 3)
n_classes=10
optim = Adam(learning_rate=0.001)
incep_model = create_model_incep(input_shape, n_classes, optim, fine_tune = 0)

```

- K-fold cross validation technique is used for evaluation. `Earlystopping` is used to optimize training time of the model.

In [12]:

```

# Train the incep model
# ModelCheckpoint callback is used to save a model or weights (in a checkpoint file) at
# EarlyStopping stops training when a monitored metric has stopped improving.
# Batch size defines the number of samples to work through before updating the internal
# The number of epochs defines the number times that the Learning algorithm will work
start = time.time()

incep_checkpoint = ModelCheckpoint(filepath='../output/incep.weights.best.hdf5', save_weights_only=True)
early_stop = EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True)
plot_loss = PlotLossesCallback()

n_split=5

for train_index,test_index in KFold(n_split).split(imgs_incep_train):

```

```

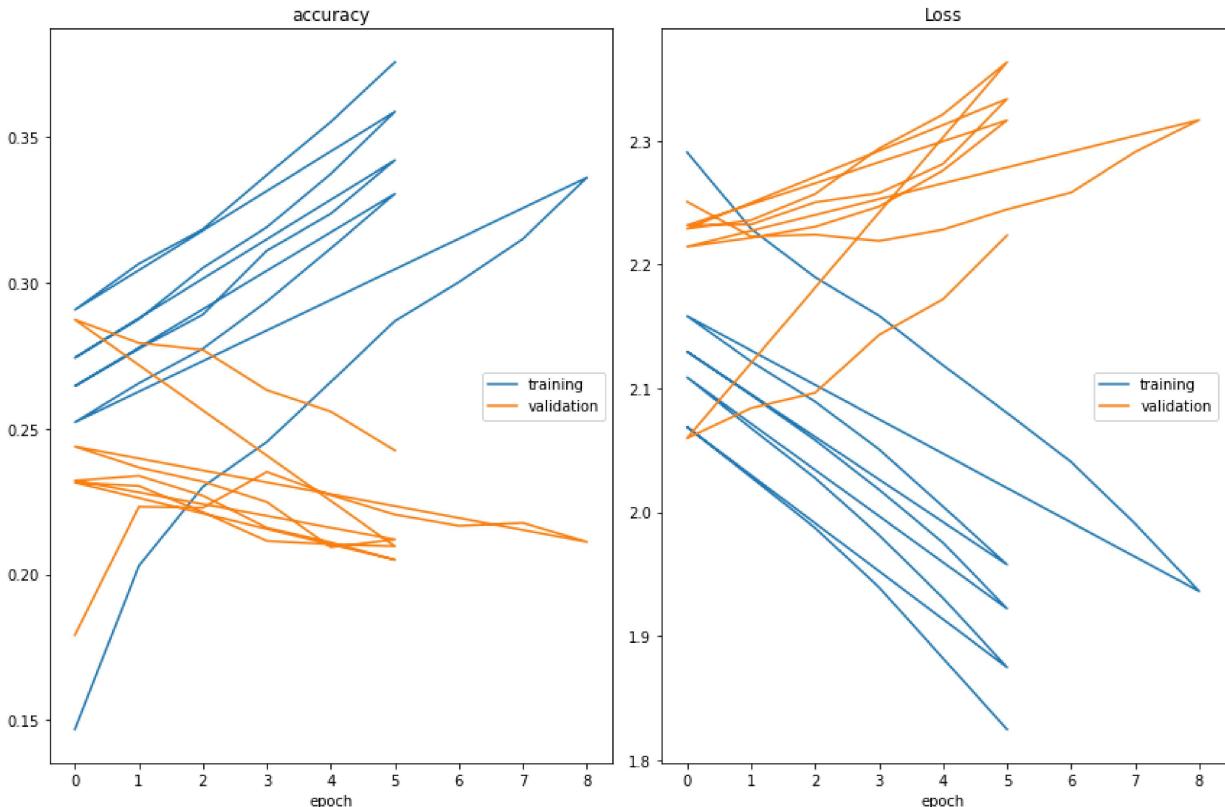
x_train,x_test=imgs_incep_train[train_index],imgs_incep_train[test_index]
y_train,y_test=noisy_labels_train[train_index],noisy_labels_train[test_index]

incep_fit = incep_model.fit(x_train,
                             tf.one_hot(y_train,10),
                             batch_size=128, # Mini-batch gradient descent
                             epochs=20,# 20
                             validation_split=0.2,
                             callbacks=[incep_checkpoint, early_stop, plot_loss],
                             verbose=1)

print('Model evaluation ',incep_model.evaluate(x_test,tf.one_hot(y_test,10)))

print('Model Training Time:', round((time.time() - start)/60, 2), 'mins')

```



```

accuracy
    training          (min:  0.147, max:  0.376, cur:  0.376)
    validation        (min:  0.179, max:  0.287, cur:  0.243)
Loss
    training          (min:  1.825, max:  2.291, cur:  1.825)
    validation        (min:  2.060, max:  2.363, cur:  2.224)
188/188 [=====] - 217s 1s/step - loss: 1.8247 - accuracy: 0.
3756 - val_loss: 2.2236 - val_accuracy: 0.2425
235/235 [=====] - 60s 252ms/step - loss: 2.1874 - accuracy: 0.2505
Model evaluation [2.1873598098754883, 0.2505333423614502]
Model Training Time: 141.55 mins

```

```

In [13]: # Prediction & Accuracy
start = time.time()
incep_model.load_weights('../output/incep.weights.best.hdf5')
incep_preds = incep_model.predict(imgs_incep_test)
incep_pred_classes = np.argmax(incep_preds, axis=1)
print('Model Test Time:', round((time.time() - start)/60, 2), 'mins')

```

```
incep_acc = accuracy_score(noisy_labels_test, incep_pred_classes)
print("incep Model Accuracy: {:.2f}%".format(incep_acc * 100))
```

391/391 [=====] - 147s 346ms/step
 Model Test Time: 2.76 mins
 incep Model Accuracy: 22.76%

In [14]: # save model
`incep_model.save('saved_models/incep/incep_model')`

WARNING:absl:Found untraced functions such as _jit_compiled_convolution_op, _jit_compiled_convolution_op, _jit_compiled_convolution_op, _jit_compiled_convolution_op, _jit_compiled_convolution_op while saving (showing 5 of 94). These functions will not be directly callable after loading.

INFO:tensorflow:Assets written to: saved_models/incep/incep_model/assets

INFO:tensorflow:Assets written to: saved_models/incep/incep_model/assets

In [13]: `def model_I(image):
 ...
 This is the MobileNetV3 predictive model that takes in the image and returns a lat
 ...
 #incep_model.Load_weights('incep.weights.best.hdf5')
 incep_preds = incep_model.predict([np.expand_dims(image, axis=0)], verbose=0);
 incep_pred_classes = np.argmax(incep_preds, axis=1)
 return int(incep_pred_classes[0])`

In [14]: `fig = plt.figure()

ax1 = fig.add_subplot(2,4,1)
ax1.imshow(imgs[0]/255)
ax2 = fig.add_subplot(2,4,2)
ax2.imshow(imgs[1]/255)
ax3 = fig.add_subplot(2,4,3)
ax3.imshow(imgs[2]/255)
ax4 = fig.add_subplot(2,4,4)
ax4.imshow(imgs[3]/255)
ax1 = fig.add_subplot(2,4,5)
ax1.imshow(imgs[4]/255)
ax2 = fig.add_subplot(2,4,6)
ax2.imshow(imgs[5]/255)
ax3 = fig.add_subplot(2,4,7)
ax3.imshow(imgs[6]/255)
ax4 = fig.add_subplot(2,4,8)
ax4.imshow(imgs[7]/255)

print('Clean labels:')
print(' '.join('%5s' % classes[clean_labels[j]] for j in range(8)))
print('Predicted baseline labels:')
print(' '.join('%5s' % classes[model_I(imgs_incep[j])] for j in range(8)))`

Clean labels:
 frog truck truck deer car car bird horse
 Predicted baseline labels:
 dog truck plane bird car car bird horse



2.3. Model II

- Model II Motivation: Since the labels for 40,000 images are noisy. We use the inception_v3 model weights and pass it through a fully connected network for training. This label-cleaning model is used to predict labels for the 40,000 images. We use the `new_labels` generated to train Model-I again.

```
In [10]: incep_ii = inception_v3.InceptionV3(include_top=False,
                                         weights='imagenet', input_shape=(75, 75, 3))

incep_ii.load_weights('../output/incep.weights.best.hdf5', by_name=True)

incep_ii_train_output = incep_ii.predict(imgs_incep)

1563/1563 [=====] - 120s 76ms/step
```

```
In [12]: x_train1 = incep_ii_train_output[:n_clean_noisy]
x_train2 = noisy_labels[:n_clean_noisy]
y_train = tf.one_hot(clean_labels, 10)

# write your code here...
#input1 = tf.keras.layers.Input(shape=(1, 1, 2048))
#flatten1 = tf.keras.layers.Flatten()(input1)
#input2 = tf.keras.layers.Input(shape=(1,))
#merged = tf.keras.layers.concatenate([flatten1, input2])
#dense1 = tf.keras.layers.Dense(128, activation='relu')(merged)
#dense2 = tf.keras.layers.Dense(64, activation='relu')(dense1)
#dropout1 = tf.keras.layers.Dropout(0.2)(dense2)
#dense4 = tf.keras.layers.Dense(10, activation='softmax')(dropout1)
#output1 = tf.keras.layers.add([dense4, input2])
#label_cleaning_model = tf.keras.models.Model([input1, input2], output1)

label_cleaning_input = np.concatenate((incep_ii_train_output.reshape(n_img, 2048), noi
```

```
label_cleaning_model = tf.keras.Sequential([
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(64, activation = 'relu'),
    tf.keras.layers.Dropout(0.5),
    tf.keras.layers.Dense(10, activation='softmax')
])

label_cleaning_model.compile(optimizer='adam',
                             loss='categorical_crossentropy',
                             metrics=['accuracy'])

#Label_cleaning_model.fit((x_train1,x_train2), y_train, batch_size=32, epochs=10, validation_split=0.2)

incep_checkpoint2 = ModelCheckpoint(filepath='lcn1.weights.best.hdf5', save_best_only=True)
early_stop2 = EarlyStopping(monitor='loss', patience=3, restore_best_weights=True, mode='min')
plot_loss2 = PlotLossesCallback()

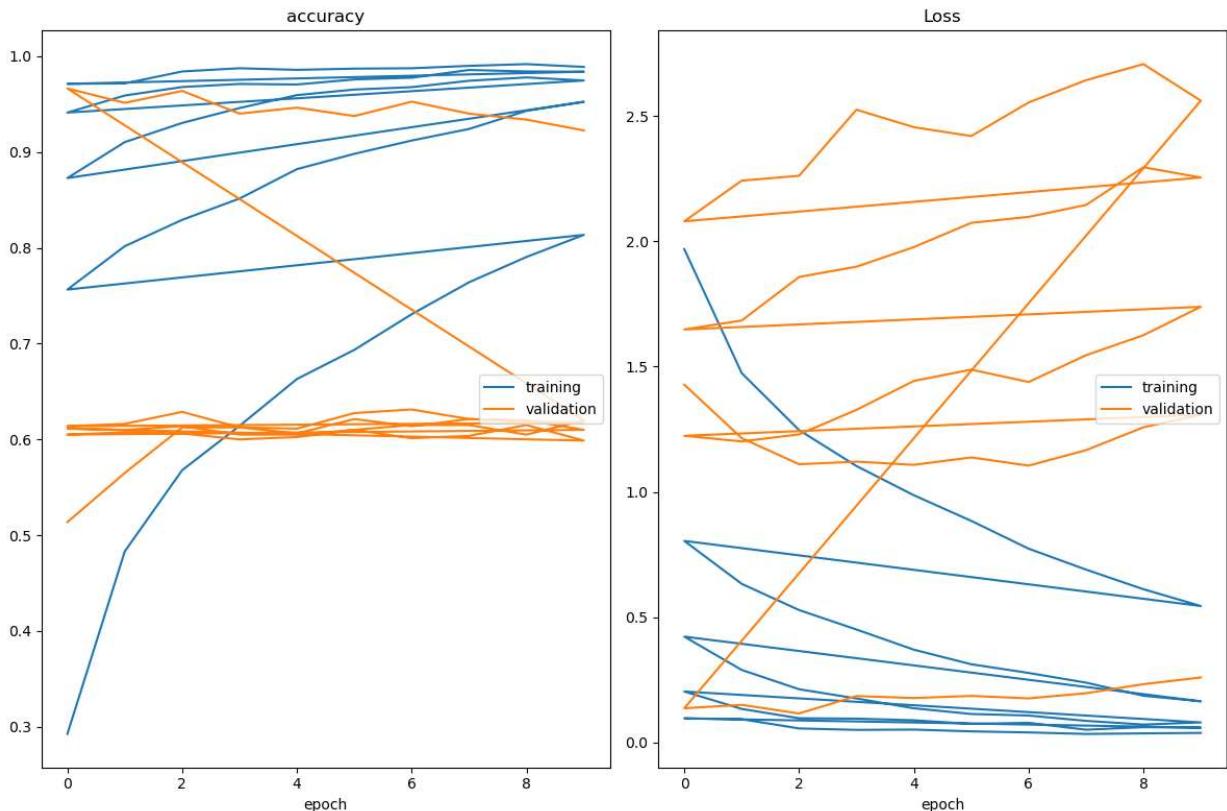
n_split=5

y_noisy1 = noisy_labels[:n_clean_noisy]

for train_index,test_index in KFold(n_split).split(label_cleaning_input[:n_clean_noisy]):
    x_train,x_test=label_cleaning_input[train_index],label_cleaning_input[test_index]
    #x_train2,x_test2=y_noisy1[train_index],y_noisy1[test_index]
    y_train,y_test=clean_labels[train_index],clean_labels[test_index]

    lcn_fit = label_cleaning_model.fit(x_train,
                                       tf.one_hot(y_train,10),
                                       batch_size=64, # Mini-batch gradient descent
                                       epochs=10,
                                       validation_split=0.1,
                                       callbacks=[incep_checkpoint2, early_stop2, plot_loss2],
                                       verbose=1)

    print('Model evaluation ',label_cleaning_model.evaluate(x_test,tf.one_hot(y_test,10))
```



```
accuracy
    training          (min:  0.293, max:  0.992, cur:  0.989)
    validation        (min:  0.514, max:  0.966, cur:  0.923)
Loss
    training          (min:  0.034, max:  1.968, cur:  0.038)
    validation        (min:  0.116, max:  2.707, cur:  0.260)
113/113 [=====] - 0s 4ms/step - loss: 0.0383 - accuracy: 0.9
887 - val_loss: 0.2596 - val_accuracy: 0.9225
63/63 [=====] - 0s 897us/step - loss: 1.2492 - accuracy: 0.8
350
Model evaluation [1.2492079734802246, 0.8349999785423279]
```

```
In [18]: # Clean Labels for all
predicted_clean_labels = label_cleaning_model.predict(label_cleaning_input[n_clean_noi])
predicted_clean_labels = predicted_clean_labels.argmax(axis=-1)
new_labels = np.append(clean_labels, predicted_clean_labels)
new_labels_oh = tf.one_hot(new_labels, depth=10)

img_class_model = tf.keras.Sequential([
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(256, activation='relu'),
    tf.keras.layers.Dense(64, activation='relu'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation='softmax')
])

img_class_model.compile(optimizer='adam',
                        loss='categorical_crossentropy',
                        metrics=['accuracy'])
```

```
1250/1250 [=====] - 1s 877us/step
```

```
In [19]: start = time.time()
incep_checkpoint1 = ModelCheckpoint(filepath='incepwlcn.weights.best.hdf5', save_best_
early_stop1 = EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True,
```

```

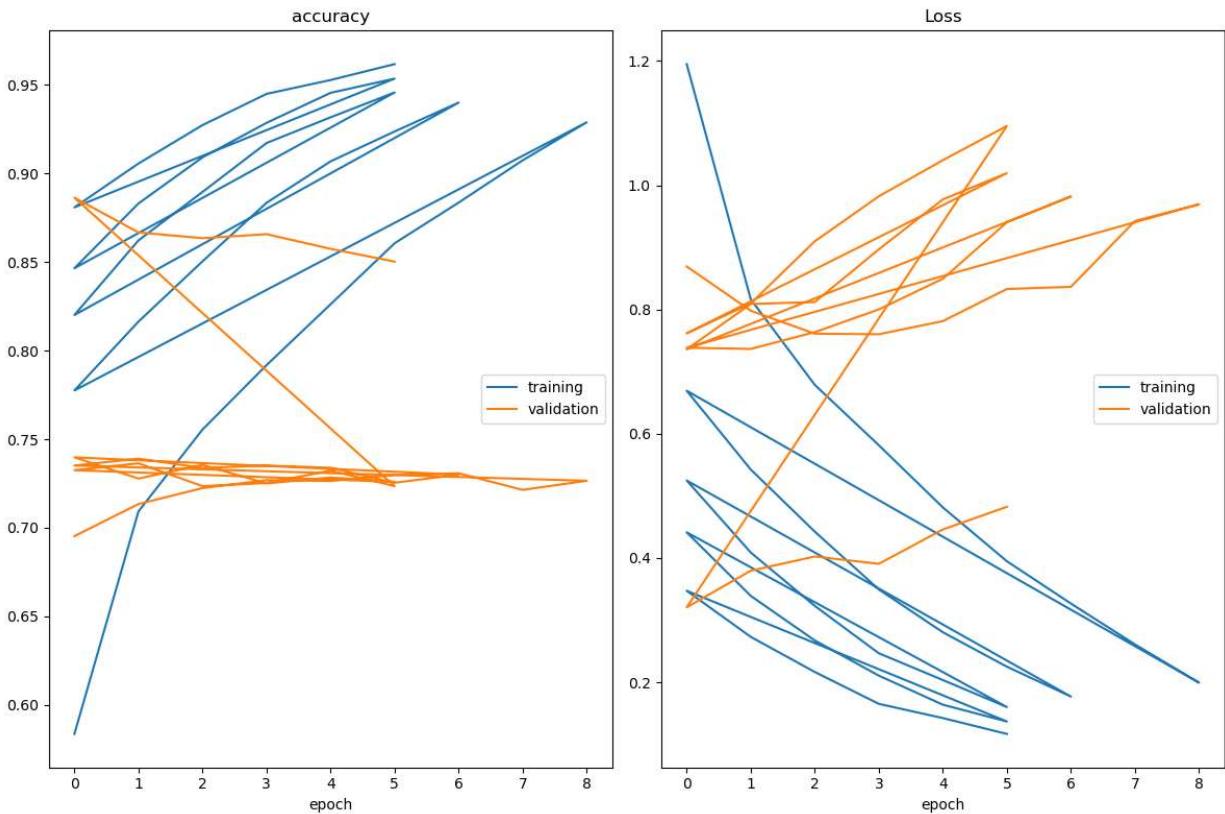
plot_loss1 = PlotLossesCallback()

for train_index,test_index in KFold(n_split).split(incep_ii_train_output):
    x_train,x_test=incep_ii_train_output[train_index],incep_ii_train_output[test_index]
    y_train,y_test=new_labels[train_index],new_labels[test_index]

    img_class_fit = img_class_model.fit(x_train,
                                         tf.one_hot(y_train,10),
                                         batch_size=128, # 64
                                         epochs=20,
                                         validation_split=0.2,
                                         callbacks=[incep_checkpoint1,early_stop1, plot_loss1],
                                         verbose=1)
    print('Model evaluation ',img_class_model.evaluate(x_test,tf.one_hot(y_test,10)))

print('Model Training Time:', round((time.time() - start)/60, 2), 'mins')

```



```

accuracy
    training          (min:  0.584, max:  0.962, cur:  0.962)
    validation        (min:  0.695, max:  0.886, cur:  0.850)

Loss
    training          (min:  0.117, max:  1.195, cur:  0.117)
    validation        (min:  0.321, max:  1.095, cur:  0.482)

250/250 [=====] - 1s 5ms/step - loss: 0.1169 - accuracy: 0.962
618 - val_loss: 0.4825 - val_accuracy: 0.8503
313/313 [=====] - 1s 2ms/step - loss: 0.7200 - accuracy: 0.758
Model evaluation [0.720008134841919, 0.7578999996185303]
Model Training Time: 0.88 mins

```

```

In [23]: # Prediction & Accuracy
start = time.time()
#incep_ii.Load_weights('incep.weights.best.hdf5',by_name=True)
#incep_ii_preds = incep_ii.predict(x_test)
img_class_model.load_weights('../output/incepwlcn.weights.best.hdf5')

```

```
incepwlcn_ii_preds = img_class_model.predict(x_test)
incep_ii_pred_classes = np.argmax(incepwlcn_ii_preds, axis=1)
print('Model Training Time:', round((time.time() - start)/60, 2), 'mins')
incep1_acc = accuracy_score(y_test, incep_ii_pred_classes)
print("Inception Model Accuracy: {:.2f}%".format(incep1_acc * 100))
```

313/313 [=====] - 0s 1ms/step

Model Training Time: 0.01 mins

Inception Model Accuracy: 67.35%

In [24]:

```
# save model
incep_ii.save('../output/saved_models/incep/incep_ii')
img_class_model.save('../output/saved_models/incep/img_class_model')
```

WARNING:tensorflow:Compiled the loaded model, but the compiled metrics have yet to be built. `model.compile_metrics` will be empty until you train or evaluate the model.

WARNING:absl:Found untraced functions such as _jit_compiled_convolution_op, _jit_compiled_convolution_op, _jit_compiled_convolution_op, _jit_compiled_convolution_op, _jit_compiled_convolution_op while saving (showing 5 of 94). These functions will not be directly callable after loading.

INFO:tensorflow:Assets written to: ../output/saved_models/incep/incep_ii\assets

INFO:tensorflow:Assets written to: ../output/saved_models/incep/incep_ii\assets

INFO:tensorflow:Assets written to: ../output/saved_models/incep/img_class_model\assets

INFO:tensorflow:Assets written to: ../output/saved_models/incep/img_class_model\assets

In [21]:

```
def model_II(image):
    ...
    This function should takes in the image of dimension 32*32*3 as input and returns
    ...
    # write your code here...
    #incep_ii.load_weights('../output/incep.weights.best.hdf5', by_name=True)
    incep_preds = incep_ii.predict([np.expand_dims(image, axis=0)], verbose=0);
    #img_class_model.load_weights('../output/incepwlcn.weights.best.hdf5')
    incep_pred_classes = np.argmax(img_class_model.predict(incep_preds, verbose=0), axis=1)
    return int(incep_pred_classes[0])
```

In [32]:

```
fig = plt.figure()

ax1 = fig.add_subplot(2,4,1)
ax1.imshow(imgs[0]/255)
ax2 = fig.add_subplot(2,4,2)
ax2.imshow(imgs[1]/255)
ax3 = fig.add_subplot(2,4,3)
ax3.imshow(imgs[2]/255)
ax4 = fig.add_subplot(2,4,4)
ax4.imshow(imgs[3]/255)
ax1 = fig.add_subplot(2,4,5)
ax1.imshow(imgs[4]/255)
ax2 = fig.add_subplot(2,4,6)
ax2.imshow(imgs[5]/255)
ax3 = fig.add_subplot(2,4,7)
ax3.imshow(imgs[6]/255)
ax4 = fig.add_subplot(2,4,8)
ax4.imshow(imgs[7]/255)

print('Clean labels:')
print(' '.join('%5s' % classes[clean_labels[j]] for j in range(8)))
```

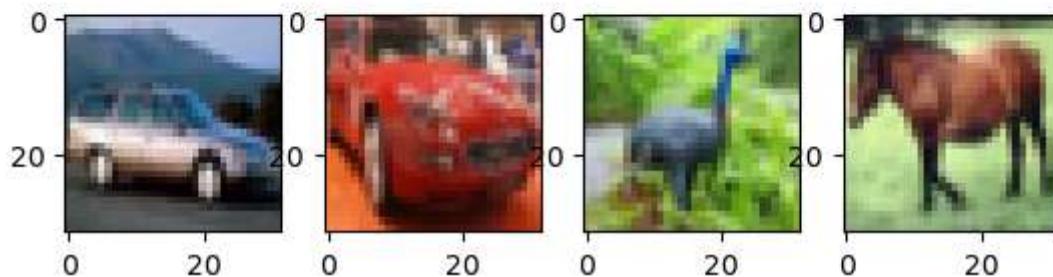
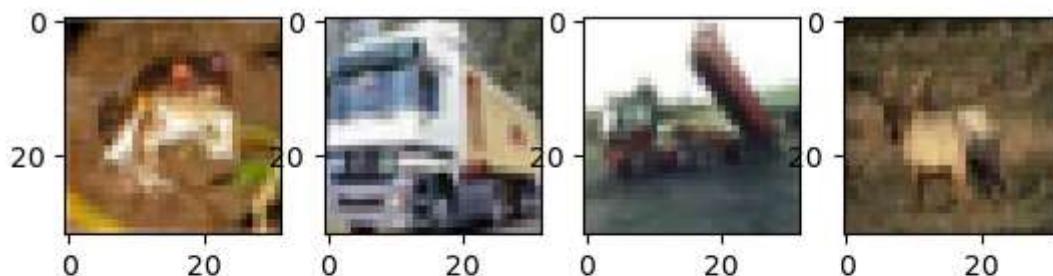
```
print('Predicted baseline labels:')
print(' '.join('%5s' % classes[model_II(imgs_incep[j])] for j in range(8)))
```

Clean labels:

frog truck truck deer car car bird horse

Predicted baseline labels:

frog truck truck deer car car bird horse



3. Evaluation

For assessment, we will evaluate your final model on a hidden test dataset with clean labels by the `evaluation` function defined as follows. Although you will not have the access to the test set, the function would be useful for the model developments. For example, you can split the small training set, using one portion for weakly supervised learning and the other for validation purpose.

```
In [33]: n_test = 10000
test_labels = np.genfromtxt('../data/test_labels.csv', delimiter=',', dtype="int8")
test_imgs = np.empty((n_test,32,32,3))
test_imgs_cnn = np.empty((n_test,75,75,3))
for i in range(n_test):
    img_fn = f'../data/test_images/{i+1:05d}.png'
    test_imgs[i,:,:,:]=cv2.cvtColor(cv2.imread(img_fn),cv2.COLOR_BGR2RGB)
    test_imgs_cnn[i,:,:,:]=cv2.resize(cv2.imread(img_fn),(75,75),interpolation=cv2.INTER_CUBIC)

test_imgs_cnn = inception_v3.preprocess_input(test_imgs_cnn)

bl_out = np.zeros(n_test)
m1_out = np.zeros(n_test)
m2_out = np.zeros(n_test)
```

```
In [26]: def model_I(image):
    '''
```

This is the MobileNetV3 predictive model that takes in the image and returns a lat

```
'''  
    incep_preds = incep_model.predict([np.expand_dims(image, axis=0)], verbose=0);  
    incep_pred_classes = np.argmax(incep_preds, axis=1)  
    return int(incep_pred_classes[0])
```

In [27]:

```
def model_II(image):  
    '''  
        This function should takes in the image of dimension 32*32*3 as input and returns  
    '''  
    # write your code here...  
    incep_preds = incep_ii.predict([np.expand_dims(image, axis=0)], verbose=0);  
    incep_pred_classes = np.argmax(img_class_model.predict(incep_preds, verbose=0), axis=1)  
    return int(incep_pred_classes[0])
```

In [28]:

```
incep_model = tf.keras.models.load_model('../output/saved_models/incep/incep_model')  
img_class_model = tf.keras.models.load_model('../output/saved_models/incep/img_class_model')  
incep_ii = inception_v3.InceptionV3(include_top=False,  
                                      weights='imagenet', input_shape=(75,75,3))  
incep_ii.load_weights('../output/incep.weights.best.hdf5', by_name=True)
```

In [40]:

```
# [DO NOT MODIFY THIS CELL]  
def evaluation(model, test_labels, test_imgs):  
    y_true = test_labels  
    y_pred = []  
    for image in test_imgs:  
        y_pred.append(model(image))  
    print(classification_report(y_true, y_pred))  
    return [item for sublist in y_pred for item in sublist]
```

In [41]:

```
# [DO NOT MODIFY THIS CELL]  
# This is the code for evaluating the prediction performance on a testset  
# You will get an error if running this cell, as you do not have the testset  
# Nonetheless, you can create your own validation set to run the evaluation  
start = time.time()  
bl_out = evaluation(baseline_model, test_labels, test_imgs)  
print('Model Evaluation Time:', round((time.time() - start)/60, 2), 'mins')
```

	precision	recall	f1-score	support
0	0.32	0.43	0.37	1005
1	0.18	0.29	0.22	974
2	0.22	0.04	0.07	1032
3	0.19	0.12	0.14	1016
4	0.24	0.48	0.32	999
5	0.22	0.13	0.16	937
6	0.26	0.35	0.30	1030
7	0.29	0.04	0.07	1001
8	0.28	0.43	0.34	1025
9	0.19	0.11	0.14	981
accuracy			0.24	10000
macro avg	0.24	0.24	0.21	10000
weighted avg	0.24	0.24	0.21	10000

Model Evaluation Time: 0.04 mins

The overall accuracy is 0.24, which is better than random guess (which should have an accuracy around 0.10). For the project, you should try to improve the performance by the following

strategies:

- Consider a better choice of model architectures, hyperparameters, or training scheme for the predictive model;
- Use both `clean_noisy_trainset` and `noisy_trainset` for model training via **weakly supervised learning** methods. One possible solution is to train a "label-correction" model using the former, correct the labels in the latter, and train the final predictive model using the corrected dataset.
- Apply techniques such as k -fold cross validation to avoid overfitting;
- Any other reasonable strategies.

In [22]:

```
# [DO NOT MODIFY THIS CELL]
# This is the code for evaluating the prediction performance on a testset
# You will get an error if running this cell, as you do not have the testset
# Nonetheless, you can create your own validation set to run the evaluation
start = time.time()
m1_out = evaluation(model_I, test_labels, test_imgs_cnn)
print('Model Evaluation Time:', round((time.time() - start)/60, 2), 'mins')
```

```

-----  

KeyboardInterrupt                                     Traceback (most recent call last)  

Input In [22], in <cell line: 7>()  

    1 # [DO NOT MODIFY THIS CELL]  

    2 # This is the code for evaluating the prediction performance on a testset  

    3 # You will get an error if running this cell, as you do not have the testset  

    4 # Nonetheless, you can create your own validation set to run the evaluation  

    5 #incep_model.load_weights('incep.weights.best.hdf5')  

    6 start = time.time()  

----> 7 m1_out = evaluation(model_I, test_labels, test_imgs_cnn)  

     8 print('Model Evaluation Time:', round((time.time() - start)/60, 2), 'mins')  

Input In [20], in evaluation(model, test_labels, test_imgs)  

    4 y_pred = []  

    5 for image in test_imgs:  

----> 6     y_pred.append(model(image))  

    7 print(classification_report(y_true, y_pred))  

    8 return y_pred  

Input In [11], in model_I(image)  

    2 '''  

    3 This is the MobileNetV3 predictive model that takes in the image and returns  

a label prediction  

    4 '''  

    5 #incep_model.load_weights('incep.weights.best.hdf5')  

----> 6 incep_preds = incep_model.predict([np.expand_dims(image, axis=0)],verbose=0);  

    7 incep_pred_classes = np.argmax(incep_preds, axis=1)  

    8 return int(incep_pred_classes[0])  

File ~\anaconda3\lib\site-packages\keras\utils\traceback_utils.py:65, in filter_traceback.<locals>.error_handler(*args, **kwargs)  

    63 filtered_tb = None  

    64 try:  

---> 65     return fn(*args, **kwargs)
    66 except Exception as e:  

    67     filtered_tb = _process_traceback_frames(e.__traceback__)
  

File ~\anaconda3\lib\site-packages\keras\engine\training.py:2220, in Model.predict(se  
lf, x, batch_size, verbose, steps, callbacks, max_queue_size, workers, use_multiproce  
ssing)
    2211     except ValueError:
    2212         warnings.warn(
    2213             "Using Model.predict with MultiWorkerMirroredStrategy "
    2214             "or TPUStrategy and AutoShardPolicy.FILE might lead to "
    (...)  

    2217             stacklevel=2,
    2218         )
-> 2220 data_handler = data_adapter.get_data_handler(  

    2221     x=x,
    2222     batch_size=batch_size,
    2223     steps_per_epoch=steps,
    2224     initial_epoch=0,
    2225     epochs=1,
    2226     max_queue_size=max_queue_size,
    2227     workers=workers,
    2228     use_multiprocessing=use_multiprocessing,
    2229     model=self,
    2230     steps_per_execution=self._steps_per_execution,
    2231 )
    2233 # Container that configures and calls `tf.keras.Callback`s.

```

```

2234 if not isinstance(callbacks, callbacks_module.CallbackList):
    main

File ~\anaconda3\lib\site-packages\keras\engine\data_adapter.py:1582, in get_data_handler(*args, **kwargs)
    main
1580 if getattr(kwargs["model"], "_cluster_coordinator", None):
1581     return _ClusterCoordinatorDataHandler(*args, **kwargs)
-> 1582 return DataHandler(*args, **kwargs)

File ~\anaconda3\lib\site-packages\keras\engine\data_adapter.py:1262, in DataHandler.__init__(self, x, y, sample_weight, batch_size, steps_per_epoch, initial_epoch, epochs, shuffle, class_weight, max_queue_size, workers, use_multiprocessing, model, steps_per_execution, distribute)
    main
1259     self._steps_per_execution = steps_per_execution
1261 adapter_cls = select_data_adapter(x, y)
-> 1262 self.adapter = adapter_cls(
    main
1263     x,
1264     y,
1265     batch_size=batch_size,
1266     steps=steps_per_epoch,
1267     epochs=epochs - initial_epoch,
1268     sample_weights=sample_weight,
1269     shuffle=shuffle,
1270     max_queue_size=max_queue_size,
1271     workers=workers,
1272     use_multiprocessing=use_multiprocessing,
1273     distribution_strategy=tf.distribute.get_strategy(),
1274     model=model,
1275 )
1277 strategy = tf.distribute.get_strategy()
1279 self._current_step = 0

File ~\anaconda3\lib\site-packages\keras\engine\data_adapter.py:347, in TensorLikeDataAdapter.__init__(self, x, y, sample_weights, sample_weight_modes, batch_size, epochs, steps, shuffle, **kwargs)
    main
344     flat_dataset = flat_dataset.shuffle(1024).repeat(epochs)
345     return flat_dataset
--> 347 indices_dataset = indices_dataset.flat_map(slice_batch_indices)
349 dataset = self.slice_inputs(indices_dataset, inputs)
351 if shuffle == "batch":

File ~\anaconda3\lib\site-packages\tensorflow\python\data\ops\dataset_ops.py:2245, in DatasetV2.flat_map(self, map_func, name)
    main
2212 def flat_map(self, map_func, name=None):
2213     """Maps `map_func` across this dataset and flattens the result.
2214
2215     The type signature is:
2216     (...)
2217     Dataset: A `Dataset`.
2218     """
-> 2245     return FlatMapDataset(self, map_func, name=name)

File ~\anaconda3\lib\site-packages\tensorflow\python\data\ops\dataset_ops.py:5484, in FlatMapDataset.__init__(self, input_dataset, map_func, name)
    main
5482 """See `Dataset.flat_map()` for details."""
5483 self._input_dataset = input_dataset
-> 5484 self._map_func = structured_function.StructuredFunctionWrapper(
5485     map_func, self._transformation_name(), dataset=input_dataset)
5486 if not isinstance(self._map_func.output_structure, DatasetSpec):
5487     raise TypeError(
5488         "The `map_func` argument must return a `Dataset` object. Got "

```

```

5489     f"({_get_type(self._map_func.output_structure)!r}).")
File ~\anaconda3\lib\site-packages\tensorflow\python\data\ops\structured_function.py:
271, in StructuredFunctionWrapper.__init__(self, func, transformation_name, dataset,
    input_classes, input_shapes, input_types, input_structure, add_to_graph, use_legacy_
function, defun_kwargs)
264     warnings.warn(
265         "Even though the `tf.config.experimental_run_functions_eagerly` "
266         "option is set, this option does not apply to tf.data functions. "
267         "To force eager execution of tf.data functions, please use "
268         "`tf.data.experimental.enable_debug_mode()`.")
269     fn_factory = trace_tf_function(defun_kwargs)
--> 271 self._function = fn_factory()
272 # There is no graph to add in eager mode.
273 add_to_graph &= not context.executing_eagerly()

File ~\anaconda3\lib\site-packages\tensorflow\python\eager\function.py:2610, in Function.get_concrete_function(self, *args, **kwargs)
2601 def get_concrete_function(self, *args, **kwargs):
2602     """Returns a `ConcreteFunction` specialized to inputs and execution context.
2603     Args:
2604         ...
2605         or `tf.Tensor` or `tf.TensorSpec`.
2606     """
--> 2610 graph_function = self._get_concrete_function_garbage_collected(
2611     *args, **kwargs)
2612 graph_function._garbage_collector.release() # pylint: disable=protected-access
2613 return graph_function

File ~\anaconda3\lib\site-packages\tensorflow\python\eager\function.py:2576, in Function._get_concrete_function_garbage_collected(self, *args, **kwargs)
2574     args, kwargs = None, None
2575     with self._lock:
--> 2576         graph_function, _ = self._maybe_define_function(args, kwargs)
2577         seen_names = set()
2578         captured = object_identity.ObjectIdentitySet(
2579             graph_function.graph.internal_captures)

File ~\anaconda3\lib\site-packages\tensorflow\python\eager\function.py:2760, in Function._maybe_define_function(self, args, kwargs)
2758     # Only get placeholders for arguments, not captures
2759     args, kwargs = placeholder_dict["args"]
--> 2760 graph_function = self._create_graph_function(args, kwargs)
2762 graph_capture_container = graph_function.graph._capture_func_lib # pylint: disable=protected-access
2763 # Maintain the list of all captures

File ~\anaconda3\lib\site-packages\tensorflow\python\eager\function.py:2670, in Function._create_graph_function(self, args, kwargs)
2665 missing_arg_names = [
2666     "%s_%d" % (arg, i) for i, arg in enumerate(missing_arg_names)
2667 ]
2668 arg_names = base_arg_names + missing_arg_names
2669 graph_function = ConcreteFunction(
--> 2670     func_graph_module.func_graph_from_py_func(
2671         self.name,
2672         self._python_function,

```

```

2673     args,
2674     kwargs,
2675     self.input_signature,
2676     autograph=self._autograph,
2677     autograph_options=self._autograph_options,
2678     arg_names=arg_names,
2679     capture_by_value=self._capture_by_value),
2680     self._function_attributes,
2681     spec=self.function_spec,
2682     # Tell the ConcreteFunction to clean up its graph once it goes out of
2683     # scope. This is not the default behavior since it gets used in some
2684     # places (like Keras) where the FuncGraph lives longer than the
2685     # ConcreteFunction.
2686     shared_func_graph=False)
2687 return graph_function

```

```

File ~\anaconda3\lib\site-packages\tensorflow\python\framework\func_graph.py:1247, in
func_graph_from_py_func(name, python_func, args, kwargs, signature, func_graph, autog
raph, autograph_options, add_control_dependencies, arg_names, op_return_value, collec
tions, capture_by_value, acd_record_initial_resource_uses)
1244 else:
1245     _, original_func = tf_decorator.unwrap(python_func)
-> 1247 func_outputs = python_func(*func_args, **func_kwargs)
1249 # invariant: `func_outputs` contains only Tensors, CompositeTensors,
1250 # TensorArrays and `None`s.
1251 func_outputs = nest.map_structure(
1252     convert, func_outputs, expand_composites=True)

```

```

File ~\anaconda3\lib\site-packages\tensorflow\python\data\ops\structured_function.py:
248, in StructuredFunctionWrapper.__init__.<locals>.trace_tf_function.<locals>.wrappe
d_fn(*args)
242 @eager_function.defun_with_attributes(
243     input_signature=structure.get_flat_tensor_specs(
244         self._input_structure),
245     autograph=False,
246     attributes=defun_kwargs)
247 def wrapped_fn(*args): # pylint: disable=missing-docstring
--> 248     ret = wrapper_helper(*args)
249     ret = structure.to_tensor_list(self._output_structure, ret)
250     return [ops.convert_to_tensor(t) for t in ret]

```

```

File ~\anaconda3\lib\site-packages\tensorflow\python\data\ops\structured_function.py:
177, in StructuredFunctionWrapper.__init__.<locals>.wrapper_helper(*args)
175 if not _should_unpack(nested_args):
176     nested_args = (nested_args,)
--> 177 ret = autograph.tf_convert(self._func, ag_ctx)(*nested_args)
178 if _should_pack(ret):
179     ret = tuple(ret)

```

```

File ~\anaconda3\lib\site-packages\tensorflow\python\autograph\impl\api.py:689, in co
nvert.<locals>.decorator.<locals>.wrapper(*args, **kwargs)
687 try:
688     with conversion ctx:
-> 689         return converted_call(f, args, kwargs, options=options)
690 except Exception as e: # pylint:disable=broad-except
691     if hasattr(e, 'ag_error_metadata'):

```

```

File ~\anaconda3\lib\site-packages\tensorflow\python\autograph\impl\api.py:377, in co
nverted_call(f, args, kwargs, caller_fn_scope, options)
374     return _call_unconverted(f, args, kwargs, options)

```

```

376 if not options.user_requested and conversion.is_allowlisted(f):
--> 377     return _call_unconverted(f, args, kwargs, options)
379 # internal_convert_user_code is for example turned off when issuing a dynamic
380 # call conversion from generated code while in nonrecursive mode. In that
381 # case we evidently don't want to recurse, but we still have to convert
382 # things like builtins.
383 if not options.internal_convert_user_code:

File ~\anaconda3\lib\site-packages\tensorflow\python\autograph\impl\api.py:458, in __call_all_unconverted(f, args, kwargs, options, update_cache)
    455     return f.__self__.call(args, kwargs)
    457 if kwargs is not None:
--> 458     return f(*args, **kwargs)
    459 return f(*args)

File ~\anaconda3\lib\site-packages\keras\engine\data_adapter.py:336, in TensorLikeDataAdapter.__init__.<locals>.slice_batch_indices(indices)
    333 flat_dataset = tf.data.Dataset.from_tensor_slices(first_k_indices)
    334 if self._partial_batch_size:
    335     index_remainder = tf.data.Dataset.from_tensors(
--> 336         tf.slice(
    337             indices, [num_in_full_batch], [self._partial_batch_size]
    338         )
    339     )
    340     flat_dataset = flat_dataset.concatenate(index_remainder)
    342 if shuffle == "batch":
    343     # 1024 is a magic constant that has not been properly evaluated

File ~\anaconda3\lib\site-packages\tensorflow\python\util\traceback_utils.py:150, in filter_traceback.<locals>.error_handler(*args, **kwargs)
    148 filtered_tb = None
    149 try:
--> 150     return fn(*args, **kwargs)
    151 except Exception as e:
    152     filtered_tb = _process_traceback_frames(e.__traceback__)

File ~\anaconda3\lib\site-packages\tensorflow\python\util\dispatch.py:1176, in add_dispatch_support.<locals>.decorator.<locals>.op_dispatch_handler(*args, **kwargs)
    1174 # Fallback dispatch system (dispatch v1):
    1175 try:
-> 1176     return dispatch_target(*args, **kwargs)
    1177 except (TypeError, ValueError):
    1178     # Note: convert_to_eager_tensor currently raises a ValueError, not a
    1179     # TypeError, when given unexpected types. So we need to catch both.
    1180     result = dispatch(op_dispatch_handler, args, kwargs)

File ~\anaconda3\lib\site-packages\tensorflow\python\ops\array_ops.py:1163, in slice(input_, begin, size, name)
    1111 @tf_export("slice")
    1112 @dispatch.add_dispatch_support
    1113 def slice(input_, begin, size, name=None):
    1114     # pylint: disable=redefined-builtin
    1115     """Extracts a slice from a tensor.
    1116
    1117     See also `tf.strided_slice`.
    (...):
    1161     A `Tensor` the same type as `input_`.
    1162     """
--> 1163     return gen_array_ops._slice(input_, begin, size, name=name)

```

```

File ~\anaconda3\lib\site-packages\tensorflow\python\ops\gen_array_ops.py:9588, in __
lice(input, begin, size, name)
    9586     pass # Add nodes to the TensorFlow graph.
    9587 # Add nodes to the TensorFlow graph.
-> 9588     , , op, outputs = op_def_library._apply_op_helper(
    9589         "Slice", input=input, begin=begin, size=size, name=name)
    9590 _result = _outputs[:]
    9591 if _execute.must_record_gradient():

File ~\anaconda3\lib\site-packages\tensorflow\python\framework\op_def_library.py:779,
in _apply_op_helper(op_type_name, name, **keywords)
    777 with g.as_default(), ops.name_scope(name) as scope:
    778     if fallback:
--> 779         _ExtractInputsAndAttrs(op_type_name, op_def, allowed_list_attr_map,
    780                         keywords, default_type_attr_map, attrs, inputs,
    781                         input_types)
    782         _ExtractRemainingAttrs(op_type_name, op_def, keywords,
    783                         default_type_attr_map, attrs)
    784     _ExtractAttrProto(op_type_name, op_def, attrs, attr_protos)

File ~\anaconda3\lib\site-packages\tensorflow\python\framework\op_def_library.py:552,
in _ExtractInputsAndAttrs(op_type_name, op_def, allowed_list_attr_map, keywords, defa
ult_type_attr_map, attrs, inputs, input_types)
    546         values = ops.convert_to_tensor(
    547             values,
    548             name=input_arg.name,
    549             as_ref=input_arg.is_ref,
    550             preferred_dtype=default_dtype)
    551     else:
--> 552         values = ops.convert_to_tensor(
    553             values,
    554             name=input_arg.name,
    555             dtype=dtype,
    556             as_ref=input_arg.is_ref,
    557             preferred_dtype=default_dtype)
558 except TypeError as err:
559     if dtype is None:

File ~\anaconda3\lib\site-packages\tensorflow\profiler\trace.py:183, in trace_
wrapper.<locals>.inner_wrapper.<locals>.wrapped(*args, **kwargs)
    181     with Trace(trace_name, **trace_kwargs):
    182         return func(*args, **kwargs)
--> 183     return func(*args, **kwargs)

File ~\anaconda3\lib\site-packages\tensorflow\python\framework\ops.py:1638, in conver
t_to_tensor(value, dtype, name, as_ref, preferred_dtype, dtype_hint, ctx, accepted_re
sult_types)
    1629     raise RuntimeError(
    1630         _add_error_prefix(
    1631             f"Conversion function {conversion_func!r} for type "
    (...))
    1634             f"actual = {ret.dtype.base_dtype.name}",
    1635             name=name))
    1637 if ret is None:
-> 1638     ret = conversion_func(value, dtype=dtype, name=name, as_ref=as_ref)
    1640 if ret is NotImplemented:
    1641     continue

File ~\anaconda3\lib\site-packages\tensorflow\python\framework\constant_op.py:343, in
_constant_tensor_conversion_function(v, dtype, name, as_ref)

```

```

340 def _constant_tensor_conversion_function(v, dtype=None, name=None,
341                                         as_ref=False):
342     _ = as_ref
--> 343     return constant(v, dtype=dtype, name=name)

File ~\anaconda3\lib\site-packages\tensorflow\python\framework\constant_op.py:267, in
constant(value, dtype, shape, name)
    170 @tf_export("constant", v1=[])
171 def constant(value, dtype=None, shape=None, name="Const"):
172     """Creates a constant tensor from a tensor-like object.
173
174     Note: All eager `tf.Tensor` values are immutable (in contrast to
175     (...))
176         ValueError: if called on a symbolic tensor.
177     """
--> 178     return constant_impl(value, dtype, shape, name, verify_shape=False,
179                           allow_broadcast=True)

File ~\anaconda3\lib\site-packages\tensorflow\python\framework\constant_op.py:289, in
_constant_impl(value, dtype, shape, name, verify_shape, allow_broadcast)
287 dtype_value = attr_value_pb2.AttrValue(type=tensor_value.tensor.dtype)
288 attrs = {"value": tensor_value, "dtype": dtype_value}
--> 289 const_tensor = g.create_op_internal( # pylint: disable=protected-access
290     "Const", [], [dtype_value.type], attrs=attrs, name=name).outputs[0]
291 if op_callbacks.should_invoke_op_callbacks():
292     # TODO(b/147670703): Once the special-op creation code paths
293     # are unified. Remove this `if` block.
294     callback_outputs = op_callbacks.invoke_op_callbacks(
295         "Const", tuple(), attrs, (const_tensor,), op_name=name, graph=g)

File ~\anaconda3\lib\site-packages\tensorflow\python\framework\func_graph.py:735, in
FuncGraph._create_op_internal(self, op_type, inputs, dtypes, input_types, name, attrs, op_def, compute_device)
733     inp = self.capture(inp)
734     captured_inputs.append(inp)
--> 735 return super(FuncGraph, self)._create_op_internal( # pylint: disable=protect
ed-access
736     op_type, captured_inputs, dtypes, input_types, name, attrs, op_def,
737     compute_device)

File ~\anaconda3\lib\site-packages\tensorflow\python\framework\ops.py:3800, in Graph.
_create_op_internal(self, op_type, inputs, dtypes, input_types, name, attrs, op_def,
compute_device)
3797 # _create_op_helper mutates the new Operation. `_mutation_lock` ensures a
3798 # Session.run call cannot occur between creating and mutating the op.
3799 with self._mutation_lock():
-> 3800     ret = Operation(
3801         node_def,
3802         self,
3803         inputs=inputs,
3804         output_types=dtypes,
3805         control_inputs=control_inputs,
3806         input_types=input_types,
3807         original_op=self._default_original_op,
3808         op_def=op_def)
3809     self._create_op_helper(ret, compute_device=compute_device)
3810 return ret

File ~\anaconda3\lib\site-packages\tensorflow\python\framework\ops.py:2108, in Operat
ion.__init__(***failed resolving arguments***)

```

```

2105     control_input_ops.append(control_op)
2107 # Initialize c_op from node_def and other inputs
-> 2108 c_op = _create_c_op(g, node_def, inputs, control_input_ops, op_def=op_def)
2109 self._init_from_c_op(c_op=c_op, g=g)
2111 self._original_op = original_op

File ~\anaconda3\lib\site-packages\tensorflow\python\util\traceback_utils.py:150, in
filter_traceback.<locals>.error_handler(*args, **kwargs)
148     filtered_tb = None
149     try:
--> 150         return fn(*args, **kwargs)
151     except Exception as e:
152         filtered_tb = _process_traceback_frames(e.__traceback__)

File ~\anaconda3\lib\site-packages\tensorflow\python\framework\ops.py:1966, in __creat
e_c_op(graph, node_def, inputs, control_inputs, op_def, extract_traceback)
1962     pywrap_tf_session.TF_SetAttrValueProto(op_desc, compat.as_str(name),
1963                                             serialized)
1965     try:
-> 1966         c_op = pywrap_tf_session.TF_FinishOperation(op_desc)
1967     except errors.InvalidArgumentError as e:
1968         # Convert to ValueError for backwards compatibility.
1969         raise ValueError(e.message)

```

KeyboardInterrupt:

In [31]:

```

# [DO NOT MODIFY THIS CELL]
# This is the code for evaluating the prediction performance on a testset
# You will get an error if running this cell, as you do not have the testset
# Nonetheless, you can create your own validation set to run the evaluation
start = time.time()
m2_out = evaluation(model_II, test_labels, test_imgs_cnn)
print('Model Evaluation Time:', round((time.time() - start)/60, 2), 'mins')

```

	precision	recall	f1-score	support
0	0.93	0.94	0.93	1005
1	0.96	0.92	0.94	974
2	0.91	0.89	0.90	1032
3	0.92	0.83	0.88	1016
4	0.86	0.91	0.88	999
5	0.91	0.89	0.90	937
6	0.89	0.96	0.92	1030
7	0.92	0.90	0.91	1001
8	0.98	0.90	0.94	1025
9	0.86	0.96	0.91	981
accuracy			0.91	10000
macro avg	0.91	0.91	0.91	10000
weighted avg	0.91	0.91	0.91	10000

Model Evaluation Time: 22.45 mins

In [42]:

```

d = {'Baseline': bl_out, 'Model I': m1_out, 'Model II': m2_out }
label_prediction = pd.DataFrame(d)
label_prediction.to_csv('../output/label_prediction.csv', index = True)

```